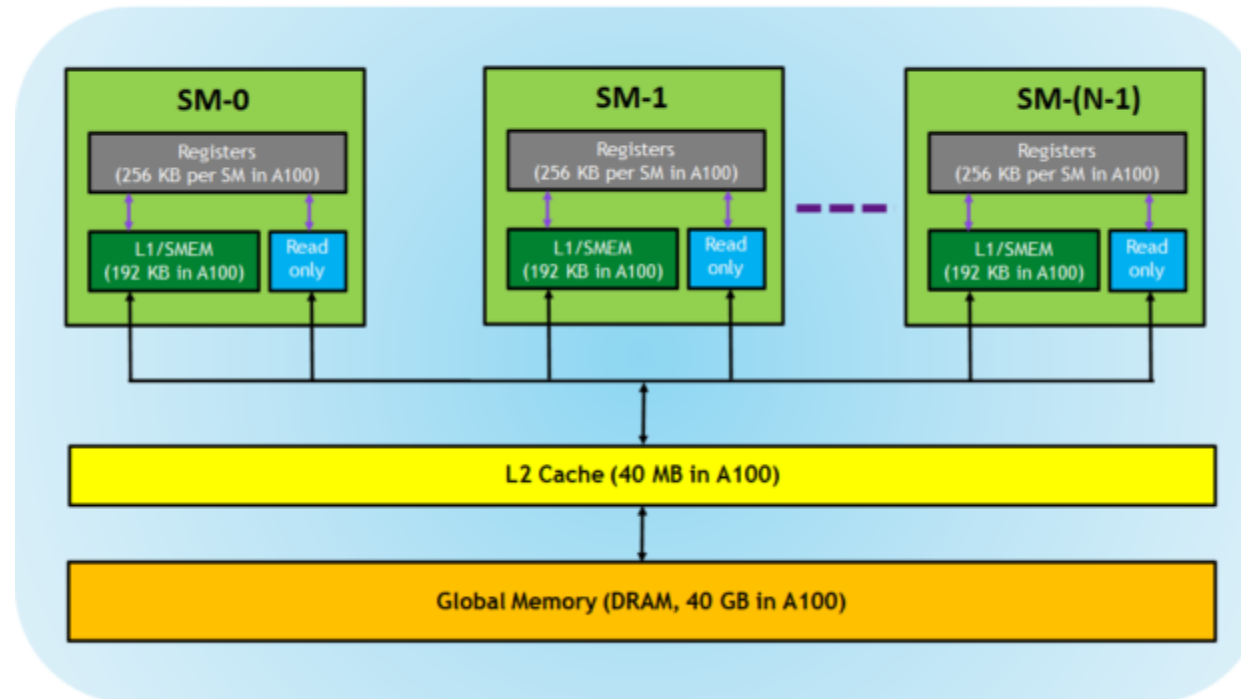


# RADL – GPU Group

- 
- 01 GEMM Optimization
  - 02 Convolution Optimazition
  - 03 CNN Benchmark
  - 04 Quantization
  - 05 Pruning
  - 06 Conclusion

# GEMM Optimization



The following memories are exposed by the GPU architecture:

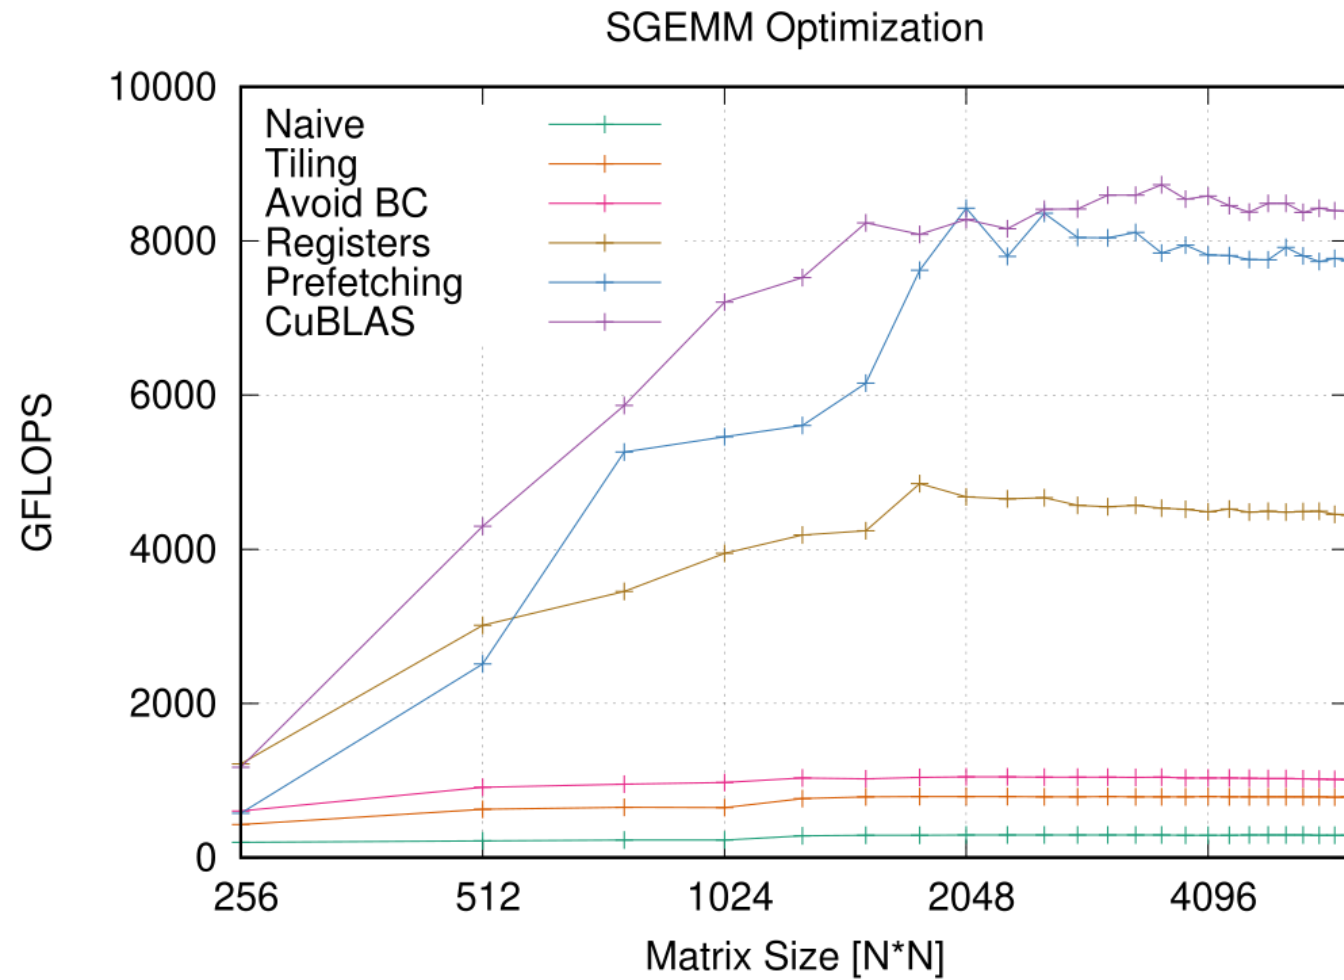
Registers; L1/Shared memory (SMEM); Read-only memory; L2 cache; Global memory

---


## Optimizing GEMM on GPU and CPU platforms share the same idea:

- Hide the memory latency with massive parallelism
- Cache-/register-level data re-use
- Manual prefetching.

The major operations in convolutional neural networks consist of matrix multiplications in both the convolutional and the fully-connected layers.



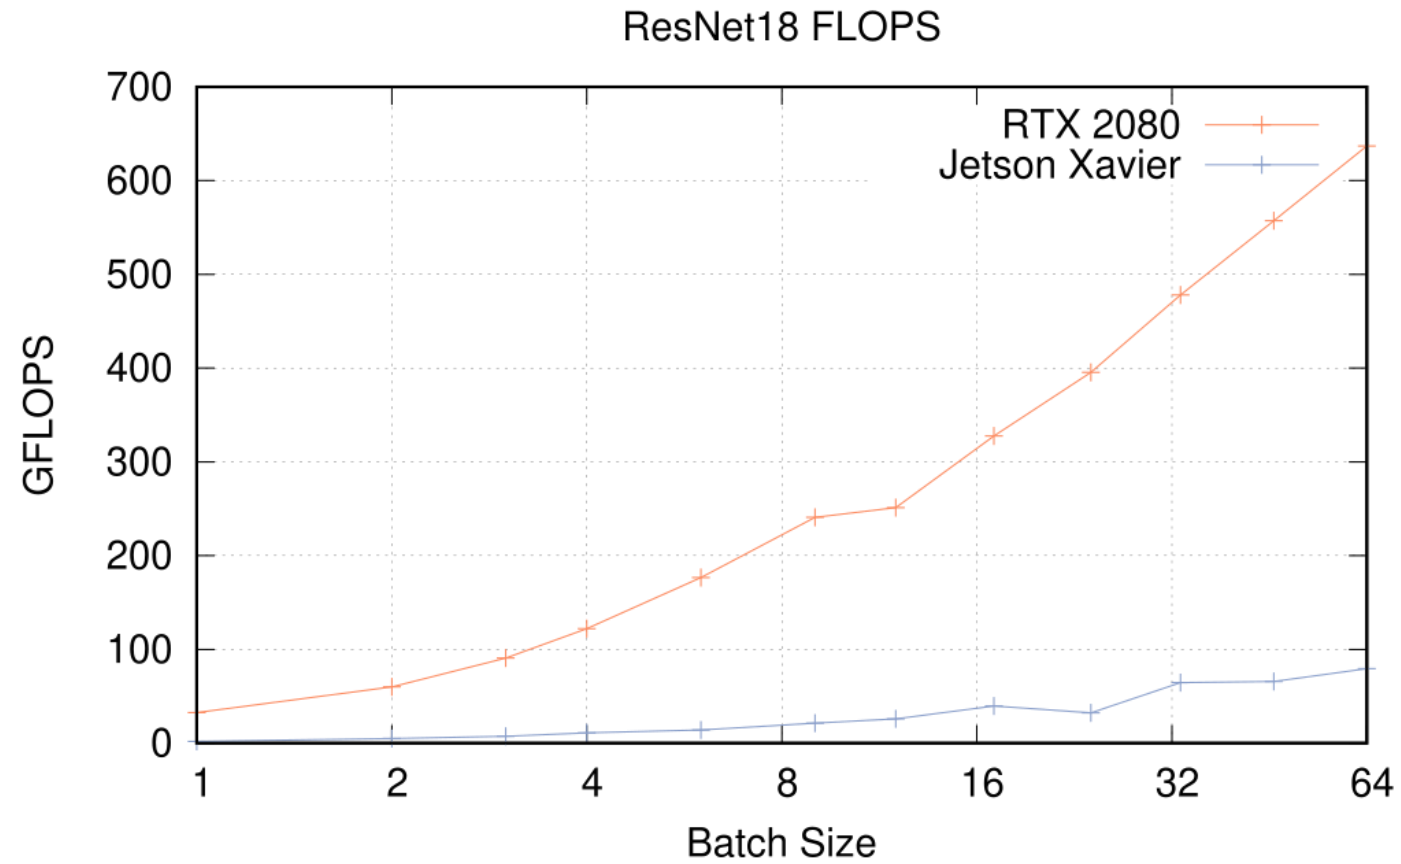
# Convolution Optimazition

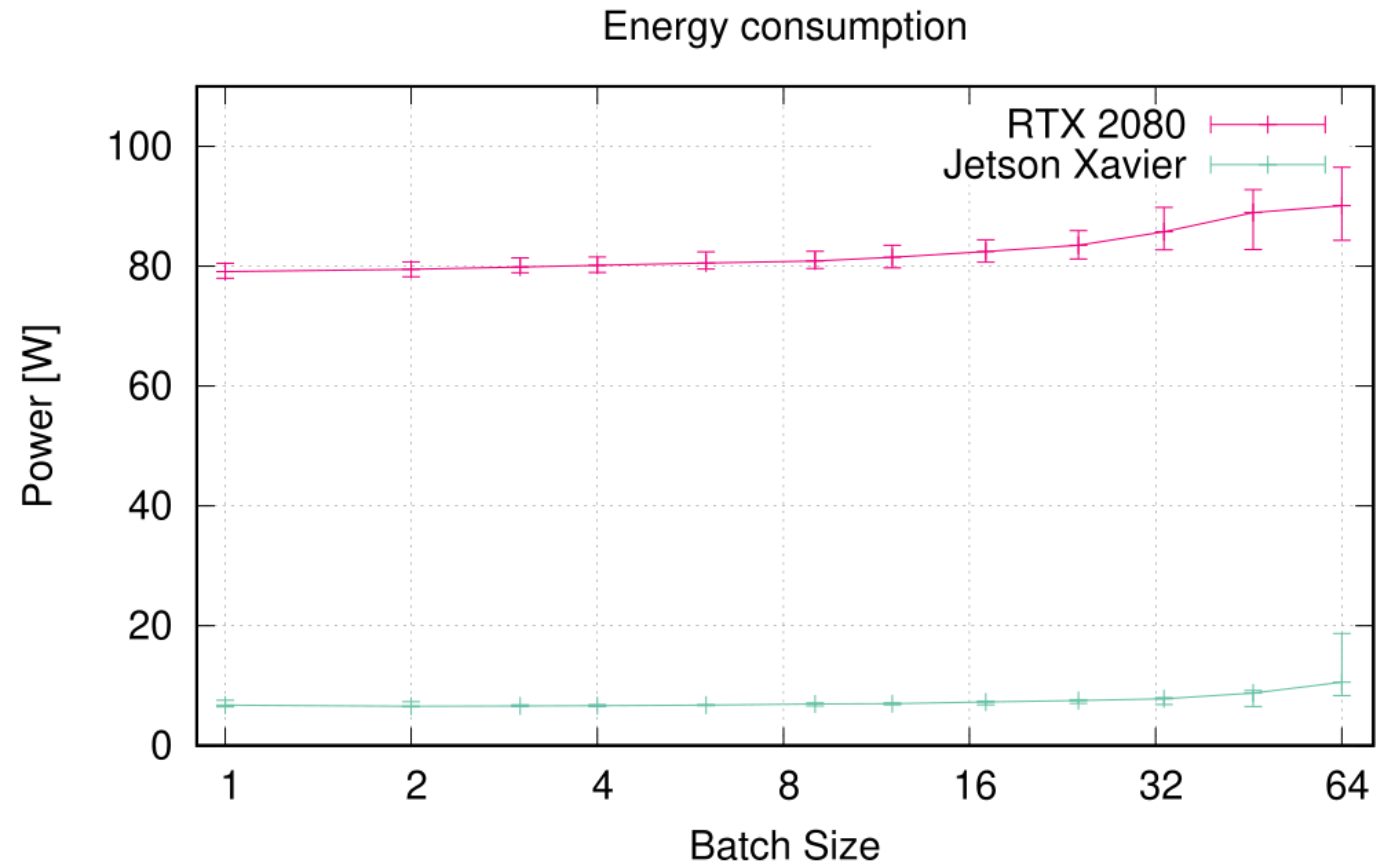




# CNN Benchmark

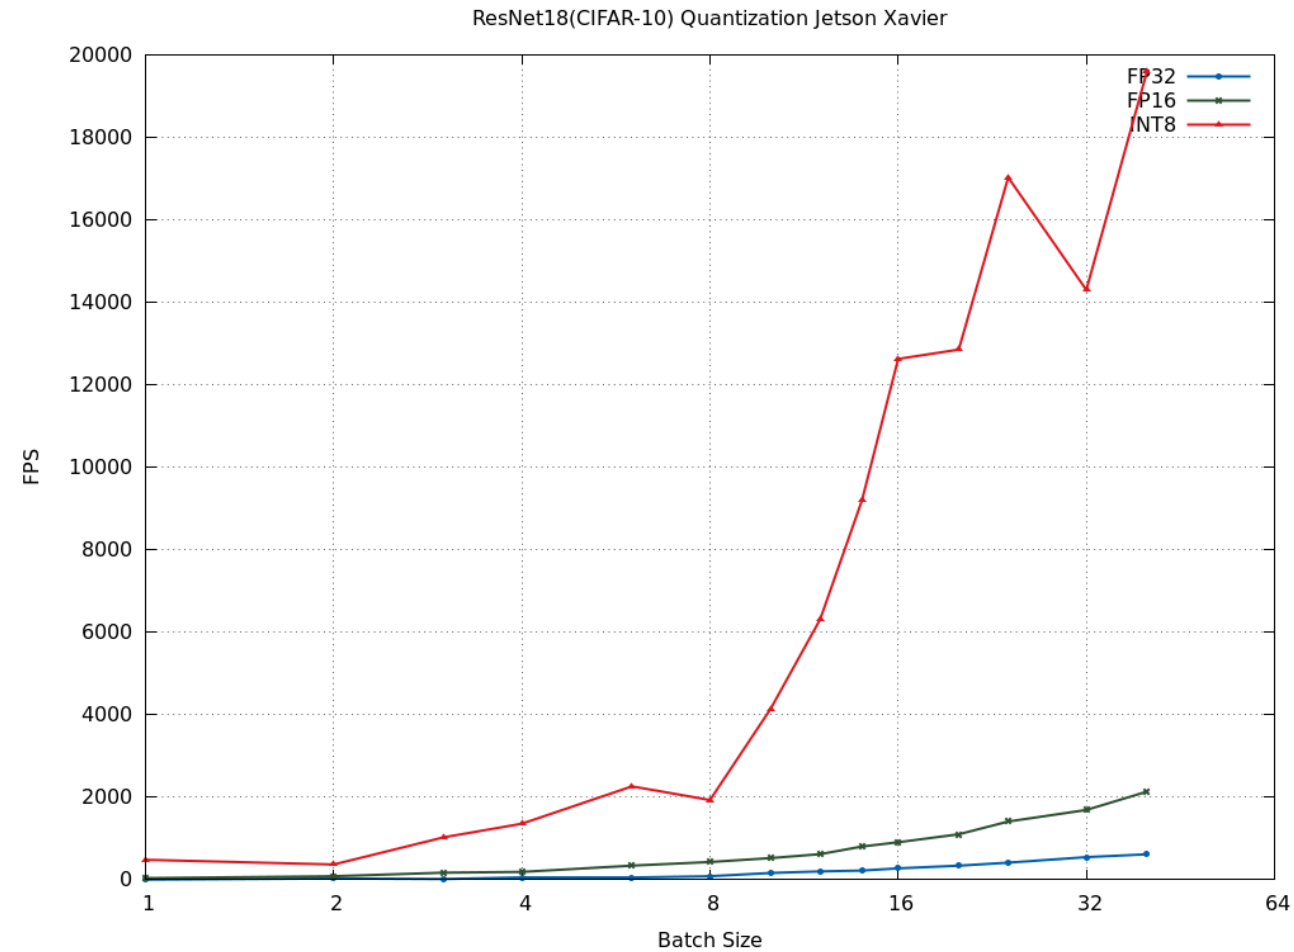
- 
- Implemented FC, Conv, ReLU, BN and a Pooling Layer using CUDA
  - Benchmarked on GPGPU(RTX 2080) and Embedding GPU(Jetson Xavier)
  - Compared the throughput and power consumption of two different platforms
  - Model: ResNet18
  - Dataset: CIFAR-10







# Quantization



# Pruning



Getting faster/smaller networks is important for running these deep learning networks on mobile devices.

The ranking can be done according to the L1/L2 norm of neuron weights, their mean activations, the number of times a neuron wasn't zero on some validation set, and other creative methods . After the pruning, the accuracy will drop

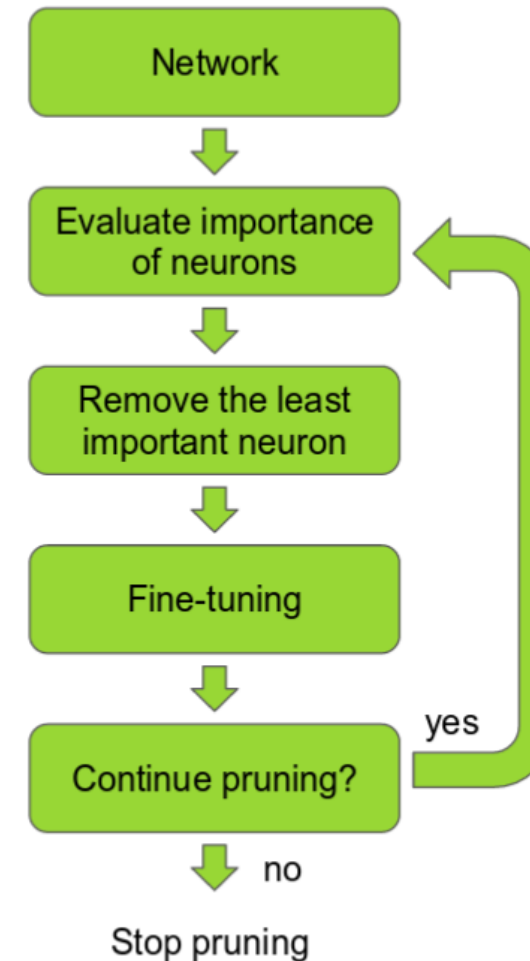
```
def apply(self, weights, amount=0.0, round_to=1) -> Sequence[int]: # return index
    if amount<=0: return []
    n = len(weights)

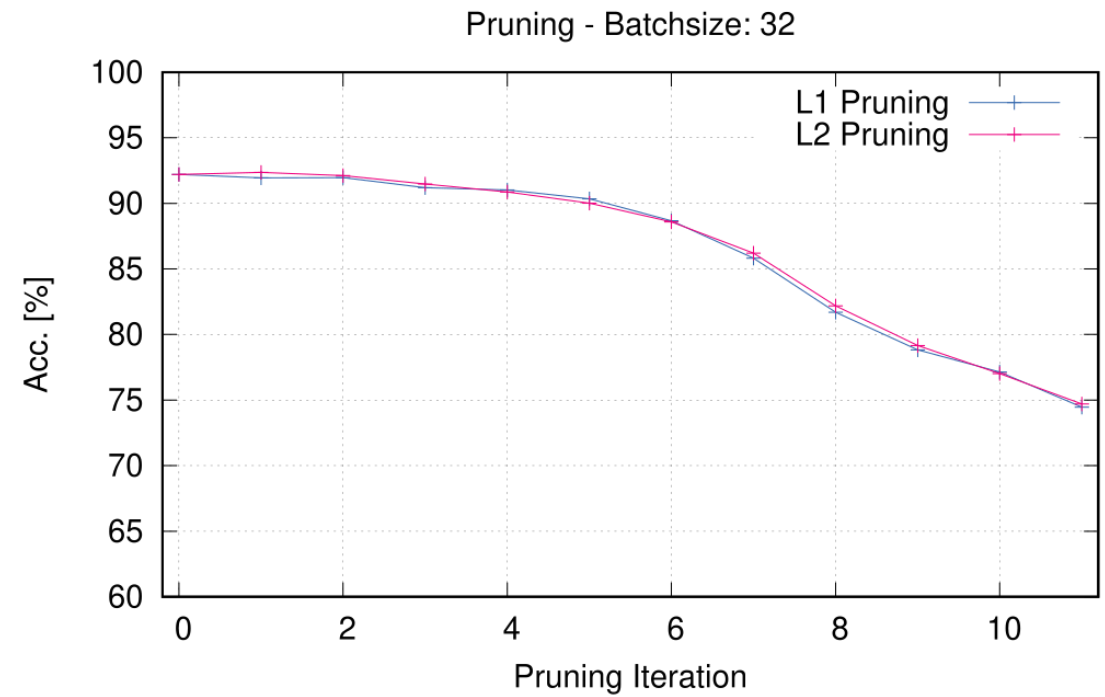
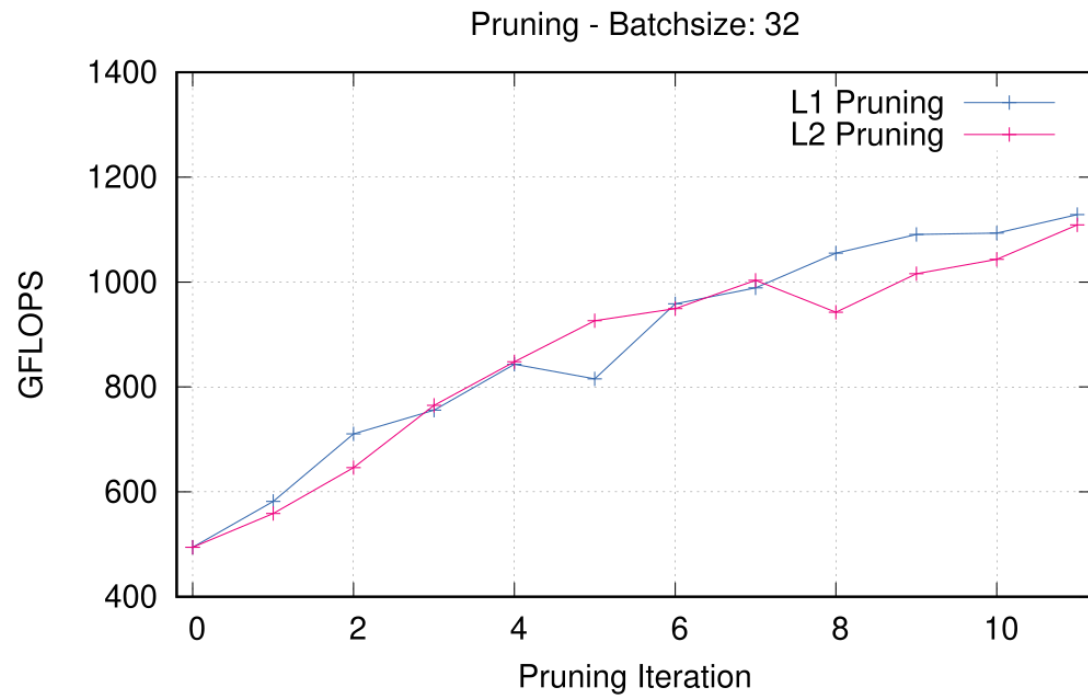
    l1_norm = torch.norm( weights.view(n, -1), p=self.p, dim=1 )

    n_to_prune = int(amount*n) if amount<1.0 else amount
    n_to_prune = round_pruning_amount(n, n_to_prune, round_to)
    if n_to_prune == 0: return []
    threshold = torch.kthvalue(l1_norm, k=n_to_prune).values
    indices = torch.nonzero(l1_norm <= threshold).view(-1).tolist()
    return indices
```


If we prune too much at once, the network might be damaged so much it won't be able to recover.

So, in practice this is an iterative process - often called 'Iterative Pruning': Prune / Train / Repeat.






Round	Params[K]	Reduced[%]	L1 Acc.[%]	L2 Acc.[%]
0	11181.642	0	92.2	92.2
1	4499.885	59.75	91.94	92.36
2	1904.148	82.97	91.96	92.11
3	850.709	92.39	91.20	91.48
4	408.764	96.34	91.03	90.87
5	210.707	98.11	90.35	90.02
6	118.955	98.93	88.68	88.59
7	70.37	99.37	85.82	86.20
8	46.003	99.58	81.70	82.19
9	32.347	99.71	78.83	79.16
10	23.606	99.78	77.14	77.02
11	17.706	99.84	74.46	74.71




L1 Norm Pruning



L2 Norm Pruning

# Conclusion



---

**Vielen Dank  
für Ihre Aufmerksamkeit!**